

# Project 1: Gym Equipment Classification

# Project 2 : Automatic Number/License Plate Detection and Recognition System

---

Sudheer kumar Tatavalu

Hima sai Kuruba

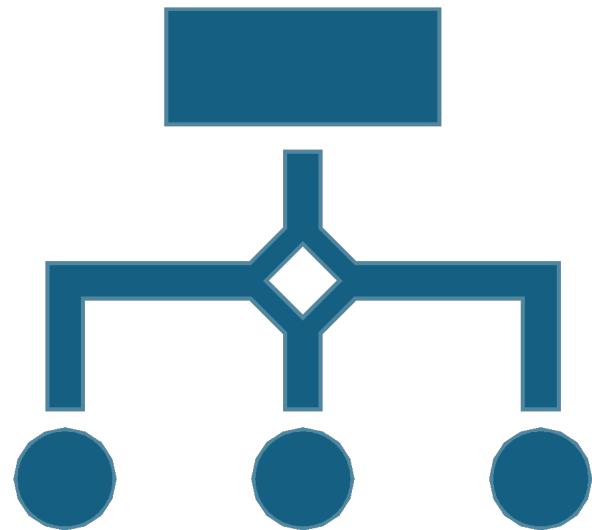
Harika Suravarapu

# Project 1: Gym Equipment Classificatio n

---



# Dataset



- Task : image classification
- Dataset Link :  
[https://drive.google.com/drive/folders/1pyTg4sFcYURECiE8vEt3KurKmDxeDkyb?usp=drive\\_link](https://drive.google.com/drive/folders/1pyTg4sFcYURECiE8vEt3KurKmDxeDkyb?usp=drive_link)
- Number of samples and image size : 1113 and 256\*256
- Number of classes : 5
- Partition: train- 80% test- 10% validation-10%
- Mean= tensor([0.7282, 0.7209, 0.7178]) and
- Standard Deviation= ([0.2826, 0.2832, 0.2816])

# Data Augmentation techniques :

- Resizing : All images are resized to a fixed size of 256×256
- Conversion to Tensor : Images are converted from PIL format to PyTorch tensors.
- Dataset Partitioning : The dataset is split into training (80%) and validation (20%) sets
- A base transformation is temporarily applied to calculate the mean and standard deviation of the dataset for normalization

# Sample images from 2 classes



Dumbbells



Rowing machine

# Architechture of NN and Loss function

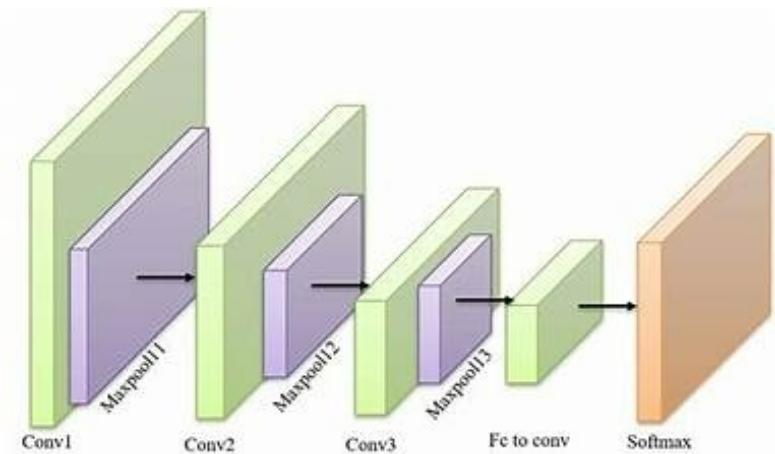


**Convolutional Neural Network (CNN)** architecture. A CNN with 3 convolutional layers, followed by pooling, fully connected layers, and dropout.



Since this is a multi-class classification problem, we used Cross-Entropy Loss function.

$$\mathcal{L} = - \sum_{c=1}^C y_c \log(\hat{y}_c)$$



# Optimization Details

## Mini-Batch Size:

- 32 (used for training, validation, and testing).

## Optimization Algorithm:

- Stochastic Gradient Descent (SGD).

## Hyperparameters for SGD:

- Learning Rate: 0.01
- Momentum: 0.9
- Weight Decay: 1e-4

## Other Tuned Hyperparameters:

- Number of Filters: 32
- Number of Convolutional Layers: 3
- Step Scheduler: StepLR (Step Size: 2, Gamma: 0)

# Hyperparameter Tuning Strategy

## Tuning Strategy:

- Manual search based on experimental results.
- Overfitting strategy on two images to validate model flexibility.

## Hyperparameter Ranges/Values:

**Learning Rate:** 0.01 to 0.001

- **Batch Size:** 32
- **Number of Layers:** 2 to 7
- **Number of Filters:** 16, 32, 64
- **Weight Decay:** 1e-4 to 5e-4
- **StepLR Gamma:** 0 to 0.1

# Hyperparameter Tuning without Learning Rate Decay

## Optimization Process Overview

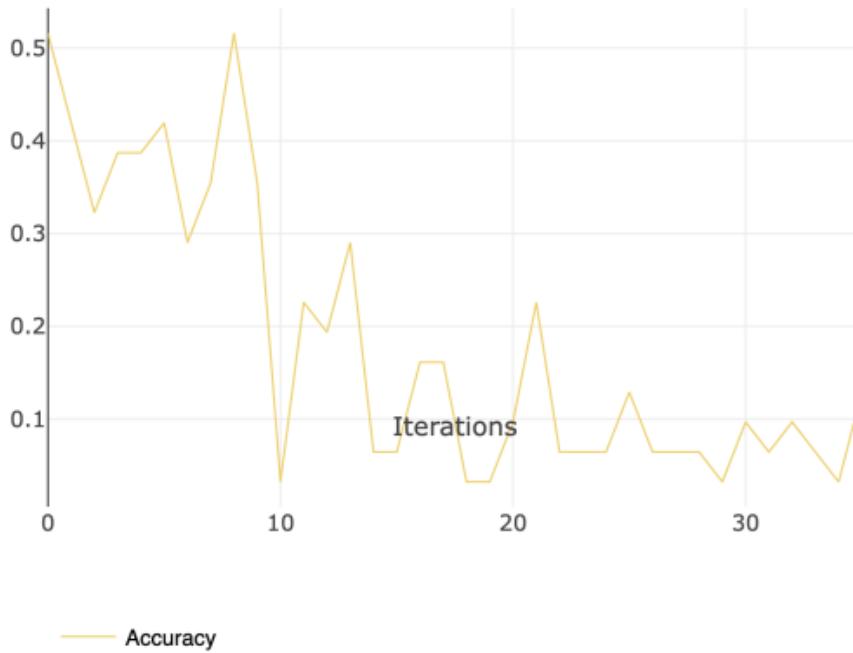
- **Optimization Method:** HyperParameterOptimizer
- **Search Strategy:** Random Search (with fallback to Bayesian or HP-Bandster if available)
- **Parameters Tuned:**
  - num\_layers (Range: 3 to 7)
  - num\_filters (Range: 5 to 32)
  - learning\_rate (LogUniform from 1e-3 to 1e-1)
  - momentum (Range: 0.7 to 0.9)
  - weight\_decay (Discrete: 1e-3, 1e-2, 1e-1)

# Hyperparameter Tuning with Learning Rate Decay

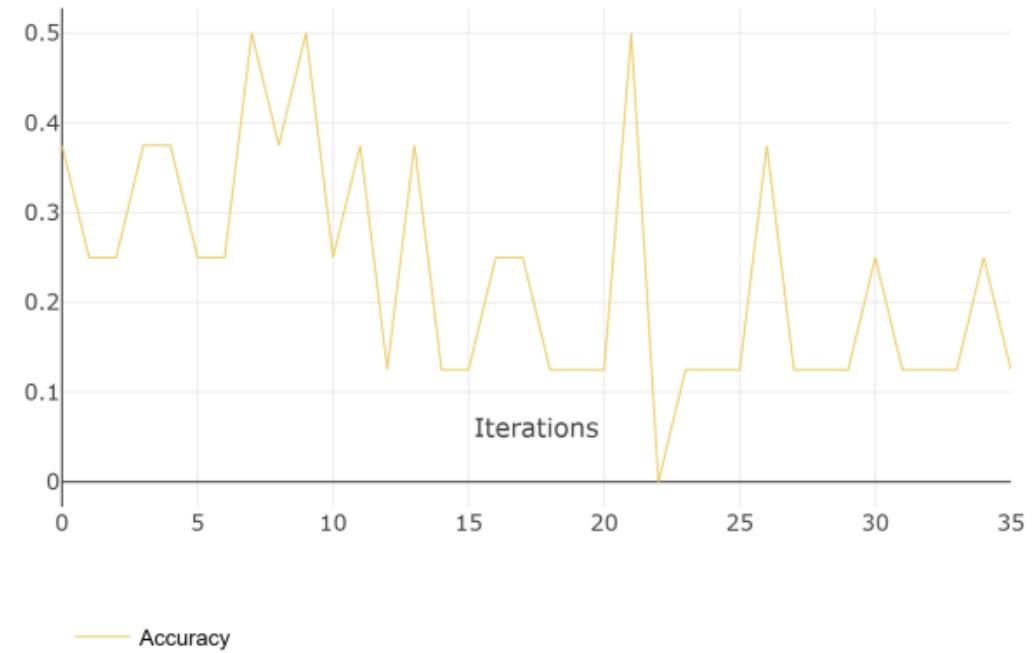
## Optimization Process Overview (with Learning Rate Decay)

- **Optimization Method:** HyperParameterOptimizer
- **Search Strategy:** Random Search (with fallback to Bayesian or HP-Bandster if available)
- **Parameters Tuned:**
  - num\_layers (Fixed: 5)
  - num\_filters (Fixed: 17)
  - learning\_rate (Fixed: 0.01)
  - momentum (Fixed: 0.74)
  - weight\_decay (Fixed: 0.1)
  - gamma (LogUniform from 1e-4 to 1e-1 for learning rate decay)

Accuracy of HPO Tuning without LR Decay



Accuracy of HPO Tuning with LR Decay



# Transfer Learning

## What is Transfer Learning?

Leverages pre-trained models (e.g., ResNet-18) on large datasets (e.g., ImageNet) for fine-tuning on specific tasks.

## Why Use Transfer Learning?

- **Faster Training:** Pre-learned features like edges and textures.
- **Improved Accuracy:** Particularly useful for small datasets.

# Experiment Setup

## Model Architecture:

**Base Model:** Pre-trained **ResNet-18** on ImageNet.

**Modifications:** Fine-tuned only the final fully connected (FC) layer for the new task (6 classes: dumbbell, cycle, etc.).

## Dataset:

**Classes:** Dumbbell and Cycle.

**Split:** 80% for training, 20% for testing.

**Image Preprocessing:** Resized to 224x224, normalized using ImageNet mean and std.

## Training Configuration:

**Batch Size:** 4

**Optimizer:** Stochastic Gradient Descent (SGD) with momentum (0.9).

**Learning Rate:** 0.001 with a learning rate scheduler (StepLR, step size=7, gamma=0.1).

**Loss Function:** Cross-Entropy Loss.

**Epochs:** 25 epochs.

**Device:** Model trained on **GPU** (if available) or **CPU**.

# Results & Analysis

## Training Accuracy and Loss:

**Accuracy:** Achieved ~85% accuracy on training set after 25 epochs.

**Loss:** Reduced to ~0.35 by the end of training.

## Test Accuracy:

**Accuracy on Test Set:** ~82% accuracy on test data.

**Confusion Matrix:** Good class separation with some misclassifications due to visual similarities.

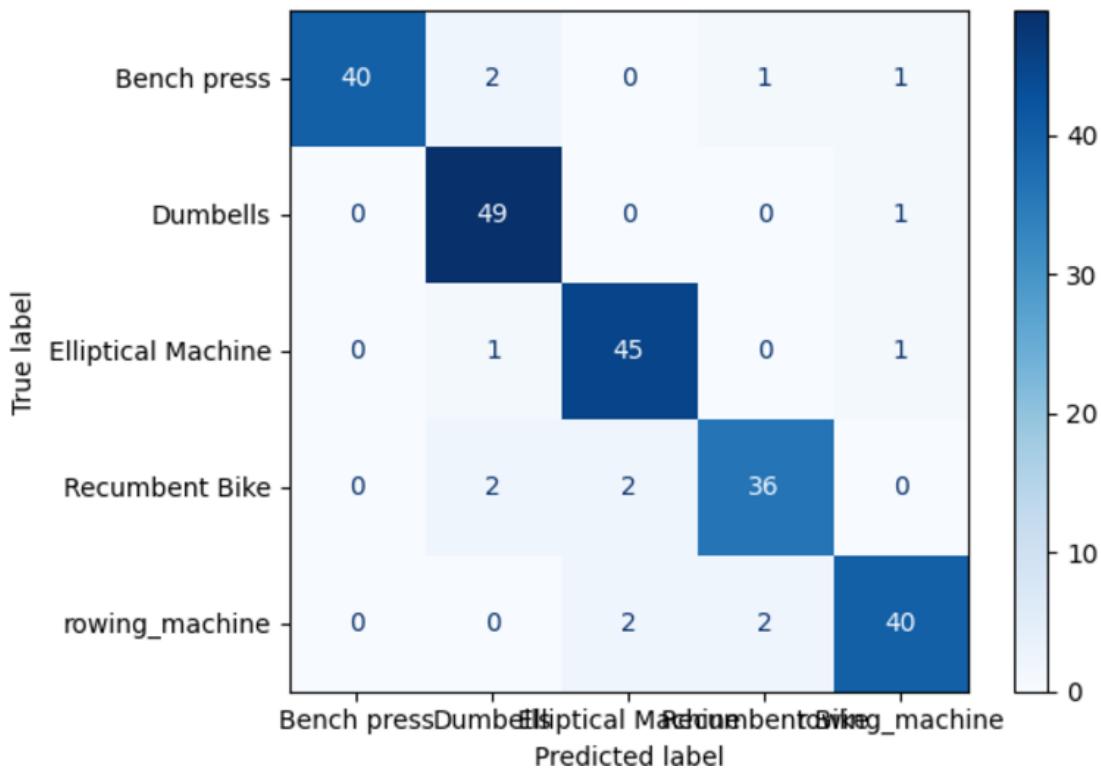
## Performance Metrics:

**Precision:** ~0.87, indicating the model is good at predicting the correct class.

**Recall:** ~0.80, showing some missed instances of the classes.

**F1-Score:** Balanced scores across classes.

# Model Evaluation



## Classification Report:

	precision	recall	f1-score	support
Bench press	1.0000	0.9091	0.9524	44
Dumbbells	0.9074	0.9800	0.9423	50
Elliptical Machine	0.9184	0.9574	0.9375	47
Recumbent Bike	0.9231	0.9000	0.9114	40
rowing_machine	0.9302	0.9091	0.9195	44
accuracy			0.9333	225
macro avg	0.9358	0.9311	0.9326	225
weighted avg	0.9351	0.9333	0.9333	225

Validation Loss: 0.2496

Validation Accuracy: 0.9333

# Transfer Learning vs Fully trained NN

Comparing the performances

Aspect	Transfer Learning (ResNet18)	Fully Trained NN (Custom CNN)
Training Time	Significantly faster due to pre-trained weights	Longer, as training is from scratch
Accuracy	Achieved high accuracy with fewer data	Higher potential accuracy with more data
Data Requirement	Effective with limited data (smaller dataset)	Needs a larger dataset to perform well
Computational Cost	Lower, efficient use of resources	Higher, more resources required
Flexibility	Less flexible, constrained by pre-trained features	More flexible, learns task-specific features

# Project 2: Automatic Number/License Plate Detection and Recognition System

---

# Data Collection and processing

Project Overview:

- **Task:** Object detection and number plate text extraction using YOLOv5 and Tesseract OCR.
- Detects three classes: **Logo**, **Number Plate**, and **Signal Lights**.

# Dataset Details:

- **Link to Dataset:** Dataset stored on [Google Drive](#) for easy access in Google Colab.
- **Number of Samples:** Approximately **300** images collected and annotated.
- **Image Size:** Images are resized to **640x640** pixels for YOLOv5 input.
- **Number of Classes: Three (3) classes:**
  - Logo
  - Number Plate
  - Signal Lights

# Data Partitioning

- **Data Splits:**  
**Training Set:** 80% of the dataset.
- **Validation Set:** 10% of the dataset.
- **Test Set:** 10% of the dataset.

- dataset/
  - └ train/
    - └ car0.jpg
    - └ car0.txt
    - └ car1.jpg
    - └ car1.txt
    - └ ...
  - └ test/
  - └ val/

# Data Normalization

## Normalization Process:

1. Images are normalized using a scale factor of **1/255**.
2. Pixel values are scaled to the range **[0, 1]** for efficient model training.
3. Mean and standard deviation normalization is handled internally by YOLOv5's pre-processing pipeline.

# Data Augmentation Techniques:

- **Augmentations Applied:**  
**Horizontal Flips:** Randomly flip images to improve robustness.
- **Random Resizing:** Vary image sizes to prevent overfitting.
- **Color Jittering:** Adjust brightness, contrast, and saturation.
- **Rotation and Scaling:** Slight rotations and rescaling to simulate varied camera angles.
- **Gaussian Noise:** Add noise for robustness against noisy inputs.

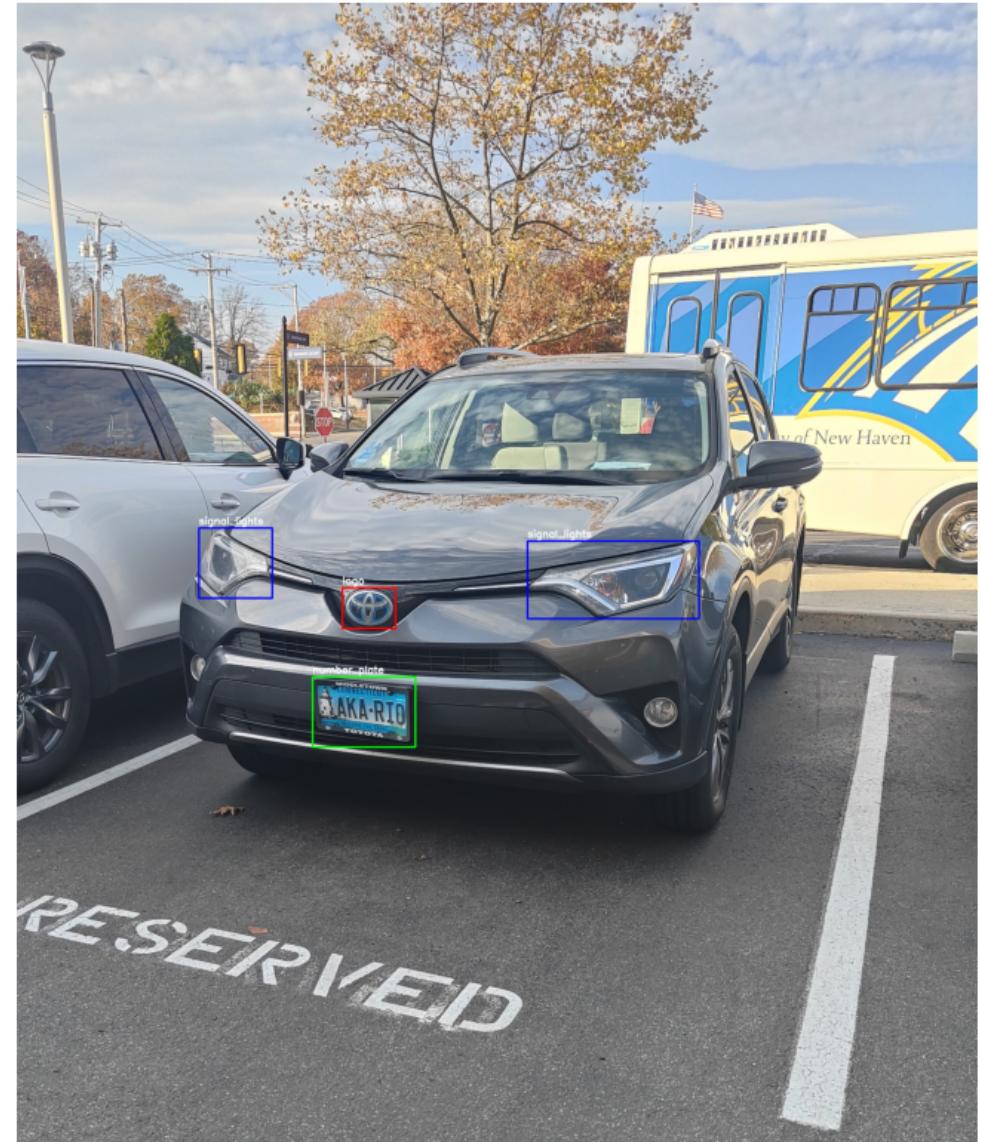
# Sample images

- Visualize two samples from different classes and their annotations
- Provide the data structure of labels for one sample

Sample1:



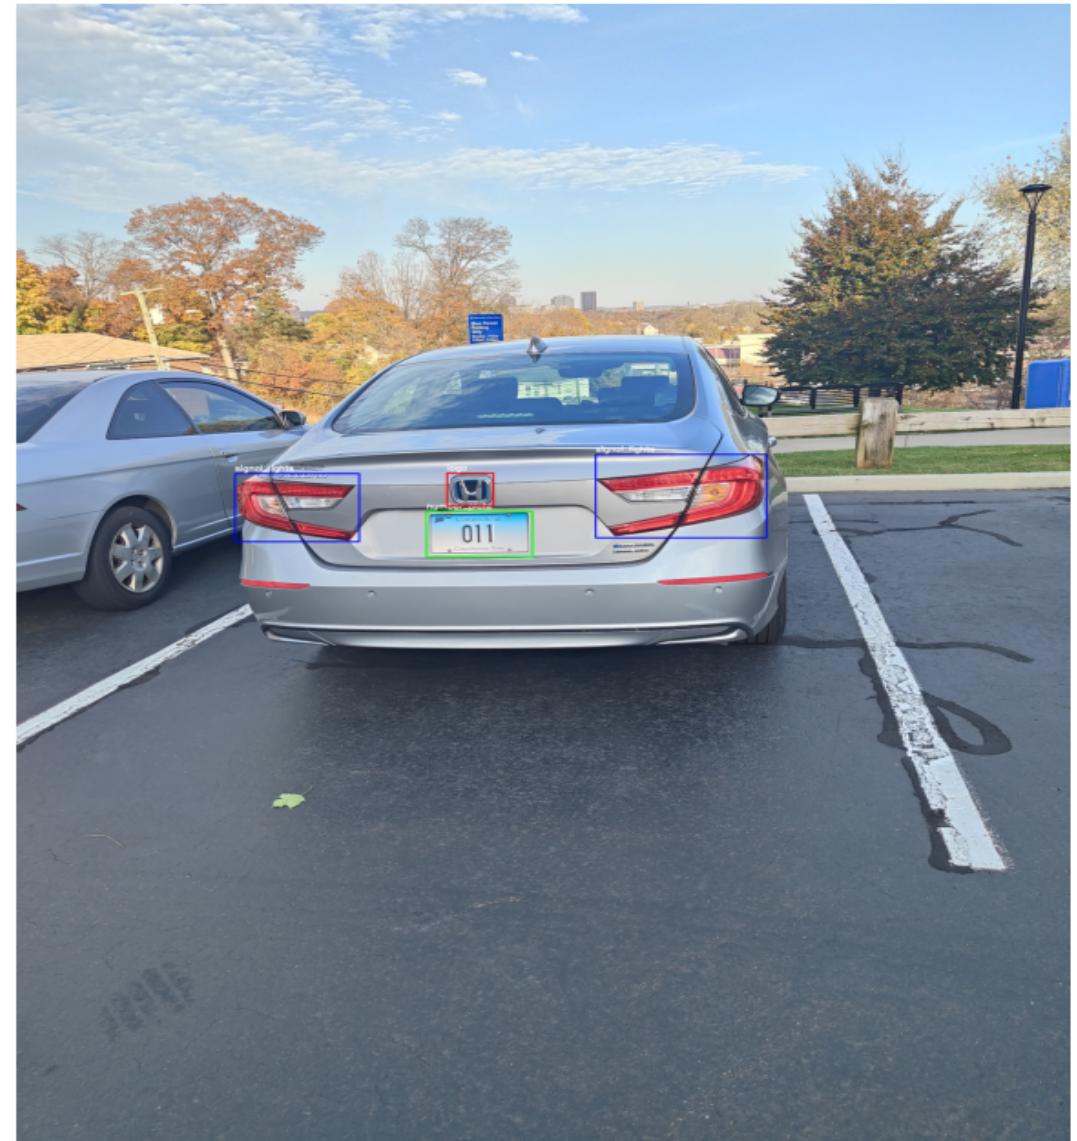
Sample1: labelled



Sample2:



Sample2: labelled



## **Sample 1: Class - Number Plate**

[Number Plate Image with Bounding Box]

- **Annotation Details:**

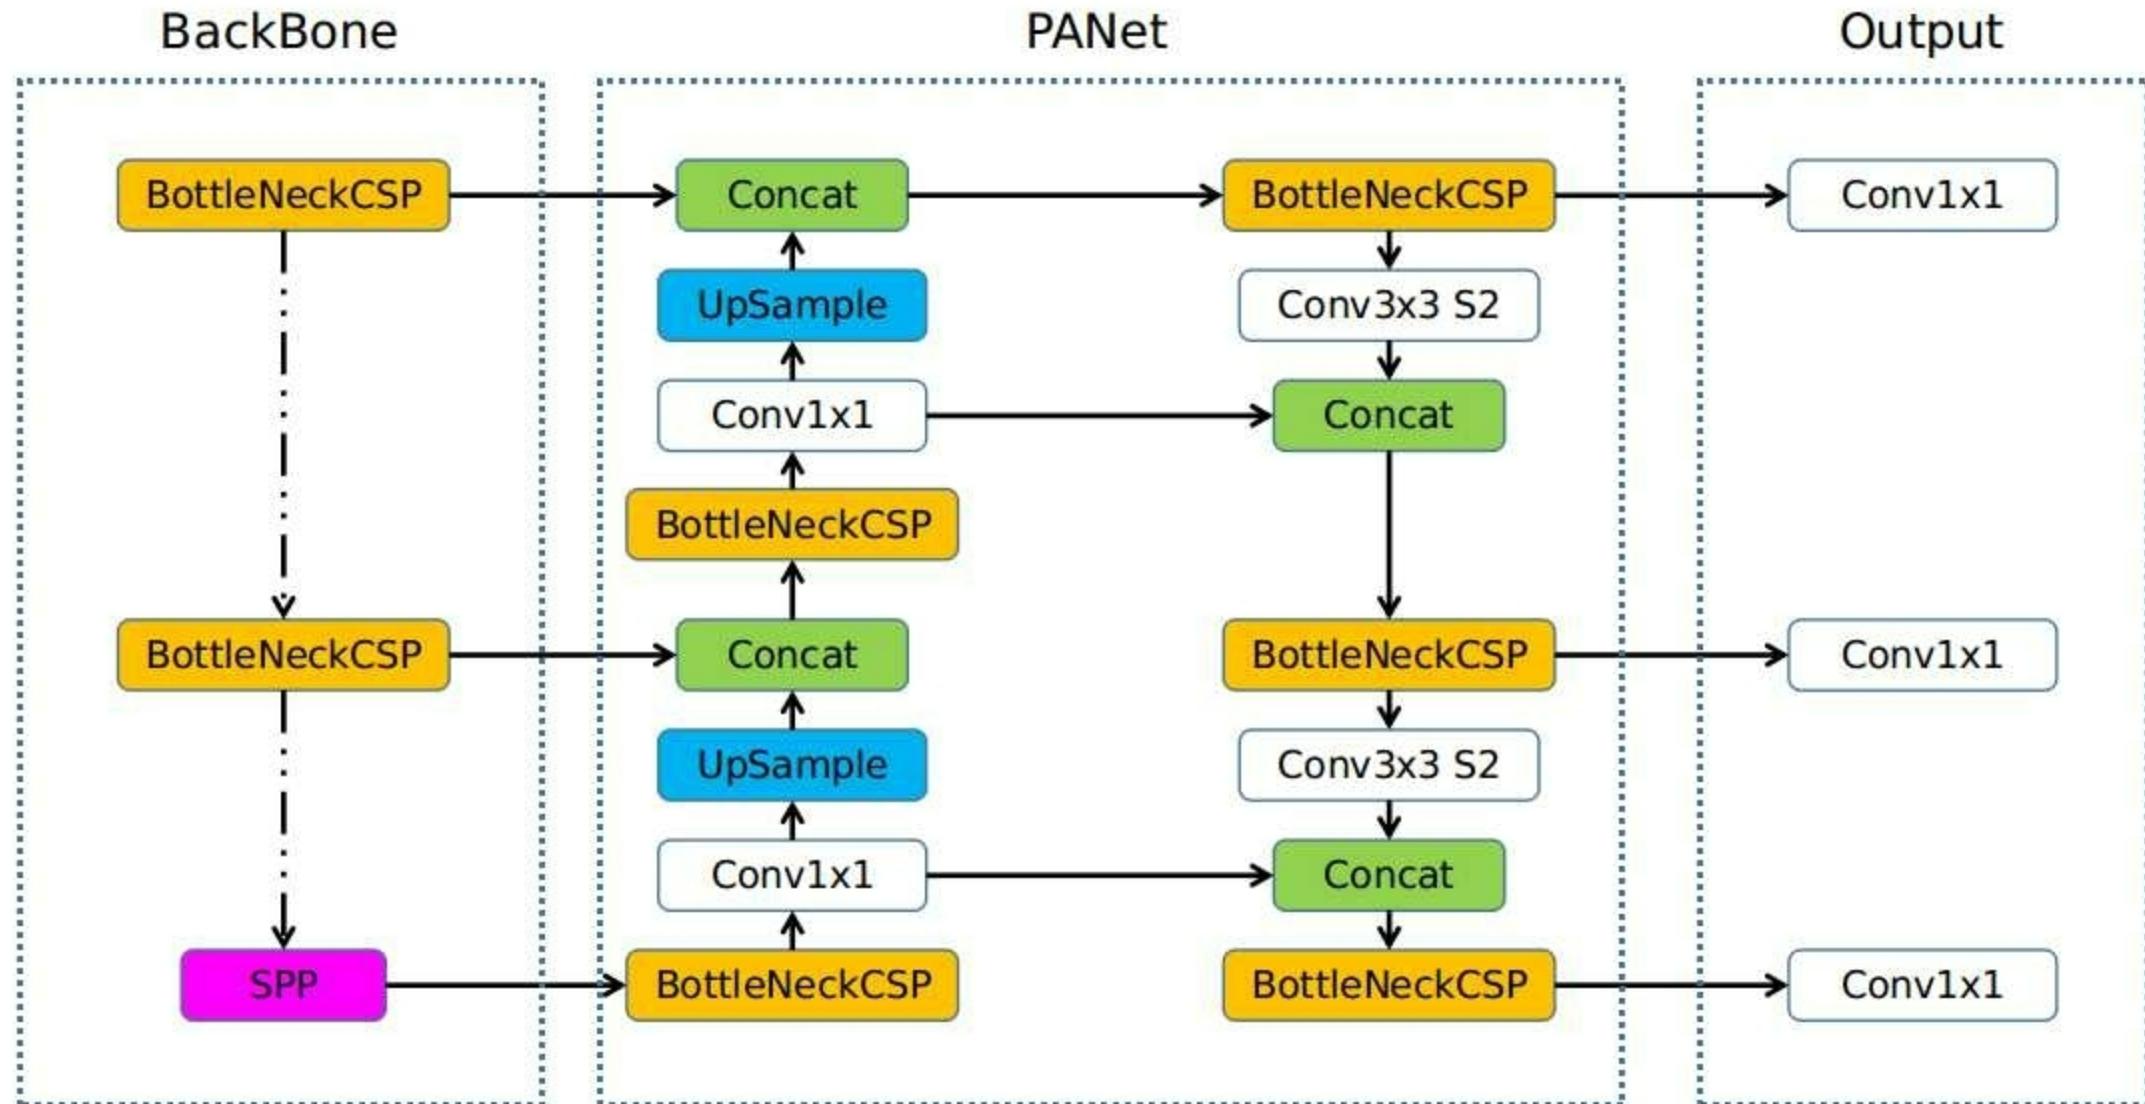
- Bounding Box: [791, 1807, 1066, 1995]  
# xmin, ymin, xmax, ymax

- Class: **Number Plate**

```
<annotation>
  <folder>Cars</folder>
  <filename>Car10.jpg</filename>
  <path>C:\Cars\Car10.jpg</path>
  <size>
    <width>2571</width>
    <height>3056</height>
    <depth>3</depth>
  </size>
  <segmented>0</segmented>
  <object>
    <name>number_plate</name>
    <pose>Unspecified</pose>
    <truncated>0</truncated>
    <difficult>0</difficult>
    <bndbox>
      <xmin>791</xmin>
      <ymin>1807</ymin>
      <xmax>1066</xmax>
      <ymax>1995</ymax>
    </bndbox>
  </object>
```

# Architecture:

# Overview of YOLOv5



## **Total Loss:**

Total Loss =  $\lambda_{\text{coord}} \cdot \text{IoU Loss} + \lambda_{\text{obj}} \cdot \text{Objectness Loss} + \lambda_{\text{cls}} \cdot \text{Classification Loss}$

Where:

$\lambda_{\text{coord}}$ ,  $\lambda_{\text{obj}}$  and  $\lambda_{\text{cls}}$  are weights for each loss component

YOLOv5 uses a combination of three loss components to optimize object detection performance:

### **1. Bounding Box Loss (IoU Loss):**

1. Measures the overlap between predicted and ground-truth boxes.
2. Formula:

$$\text{IoU} = \text{Area of Overlap} / \text{Area of Union}$$

### **1. Class Probability Loss (ClassificObjectness Loss (Confidence Loss):**

1. Penalizes incorrect predictions of whether an object exists in a particular grid cell.
2. Uses Binary Cross-Entropy (BCE) loss to predict the presence of an object.

### **2. ation Loss):**

1. Uses BCE or Cross-Entropy loss to classify objects into the correct categories.

## NN architecture:

	from	n	params	module	arguments	
0		-1	1	3520	models.common.Conv	[3, 32, 6, 2, 2]
1		-1	1	18560	models.common.Conv	[32, 64, 3, 2]
2		-1	1	18816	models.common.C3	[64, 64, 1]
3		-1	1	73984	models.common.Conv	[64, 128, 3, 2]
4		-1	2	115712	models.common.C3	[128, 128, 2]
5		-1	1	295424	models.common.Conv	[128, 256, 3, 2]
6		-1	3	625152	models.common.C3	[256, 256, 3]
7		-1	1	1180672	models.common.Conv	[256, 512, 3, 2]
8		-1	1	1182720	models.common.C3	[512, 512, 1]
9		-1	1	656896	models.common.SPPF	[512, 512, 5]
10		-1	1	131584	models.common.Conv	[512, 256, 1, 1]
11		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[ -1, 6 ]	1	0	models.common.Concat	[1]	
13		-1	1	361984	models.common.C3	[512, 256, 1, False]
14		-1	1	33024	models.common.Conv	[256, 128, 1, 1]
15		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16	[ -1, 4 ]	1	0	models.common.Concat	[1]	
17		-1	1	90880	models.common.C3	[256, 128, 1, False]
18		-1	1	147712	models.common.Conv	[128, 128, 3, 2]
19	[ -1, 14 ]	1	0	models.common.Concat	[1]	
20		-1	1	296448	models.common.C3	[256, 256, 1, False]
21		-1	1	590336	models.common.Conv	[256, 256, 3, 2]
22	[ -1, 10 ]	1	0	models.common.Concat	[1]	
23		-1	1	1182720	models.common.C3	[512, 512, 1, False]
24	[ 17, 20, 23 ]	1	21576	models.yolo.Detect	[3, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90, 156, 198, 373, 326]], [128, 256, 512]]	

YOLOv5s summary: 214 layers, 7027720 parameters, 7027720 gradients, 16.0 GFLOPs

# Transfer Learning:

```
!python train.py --data data.yaml --weights yolov5s.pt \
--img 640 --epochs 10 --batch-size 8 --name yolo_initial2 \
--freeze 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14
```

... 100% 755k/755k [00:00<00:00, 14.2MB/s]  
Overriding model.yaml nc=80 with nc=3

	from	n	params	module	arguments	
0		-1	1	3520	models.common.Conv	[3, 32, 6, 2, 2]
1		-1	1	18560	models.common.Conv	[32, 64, 3, 2]
2		-1	1	18816	models.common.C3	[64, 64, 1]
3		-1	1	73984	models.common.Conv	[64, 128, 3, 2]
4		-1	2	115712	models.common.C3	[128, 128, 2]
5		-1	1	295424	models.common.Conv	[128, 256, 3, 2]
6		-1	3	625152	models.common.C3	[256, 256, 3]
7		-1	1	1180672	models.common.Conv	[256, 512, 3, 2]
8		-1	1	1182720	models.common.C3	[512, 512, 1]
9		-1	1	656896	models.common.SPPF	[512, 512, 5]
10		-1	1	131584	models.common.Conv	[512, 256, 1, 1]
11		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[ -1, 6]	1		0	models.common.Concat	[1]
13		-1	1	361984	models.common.C3	[512, 256, 1, False]
14		-1	1	33024	models.common.Conv	[256, 128, 1, 1]
15		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16	[ -1, 4]	1		0	models.common.Concat	[1]

# Training

- Mini-batch size
- Optimization algorithm and its hyperparameter settings (e.g., learning rate, weight decay)
- Hyperparameter search

# Mini-batch Size and Optimization

## Mini-batch Size:

- Mini-batch size of **8** provides a balance between training efficiency and memory constraints.

**Optimization Algorithm: SGD with Momentum**  
**Purpose:** Helps accelerate convergence by incorporating past gradient directions.

## Hyperparameter setting:

- **Momentum** (0.7 to 1.0)
- **Learning Rate** (0.0001 to 0.01)
- **Weight Decay:** (0, 1.0e-9, 1.0e-5)

# Hyperparameter Search

## **Random Search for Hyperparameter Tuning**

**Method:** Random Search was used to explore a wide range of hyperparameters.

### **Best Hyperparameters Achieved:**

- **Learning Rate:** 0.01
- **Weight Decay:** 1e-09
- **Momentum:** 0.8151
- This combination of hyperparameters achieved the best model performance

# **Loss Comparison (Training vs. Validation)**

- Epoch 0: Train Box Loss: 0.11641      Epoch 24: Train Box Loss: 0.036988
- Train Object Loss: 0.051936
- Train Class Loss: 0.039596
- Validation Box Loss: 0.10114
- Validation Object Loss: 0.043514
- Validation Class Loss: 0.035188
- Train Object Loss: 0.029989
- Train Class Loss: 0.0044069
- Validation Box Loss: 0.037226
- Validation Object Loss: 0.020346
- Validation Class Loss: 0.0036811

## **Key Observations from Data:**

- As training progresses, we expect both **training loss** and **validation loss** to decrease.

# Performance Metrics (Precision, Recall, mAP)

- Epoch 0:**Precision**: 0.073413
- **Recall**: 0.0011638
- **mAP@0.5**: 0.00083556
- **mAP@0.5:0.95**: 0.00023351
- Epoch 24:**Precision**: 0.92911
- **Recall**: 0.0044069
- **mAP@0.5**: 0.92868
- **mAP@0.5:0.95**: 0.56218

## Key Observations:

- **Precision** and **Recall** increase significantly as training progresses, showing that the model is getting better at detecting objects.
- **mAP** (both at IoU 0.5 and 0.5:0.95) improves over time, indicating that the model is improving its detection quality across both easy and harder detection cases

# Pre- and Post-Training Comparison (Summary)

## Training Loss:

- Training loss has significantly decreased, indicating that the model has learned the patterns in the data. For example:
  - **Epoch 0: Train Box Loss** = 0.11641, **Train Object Loss** = 0.051936
  - **Epoch 24: Train Box Loss** = 0.036988, **Train Object Loss** = 0.029989

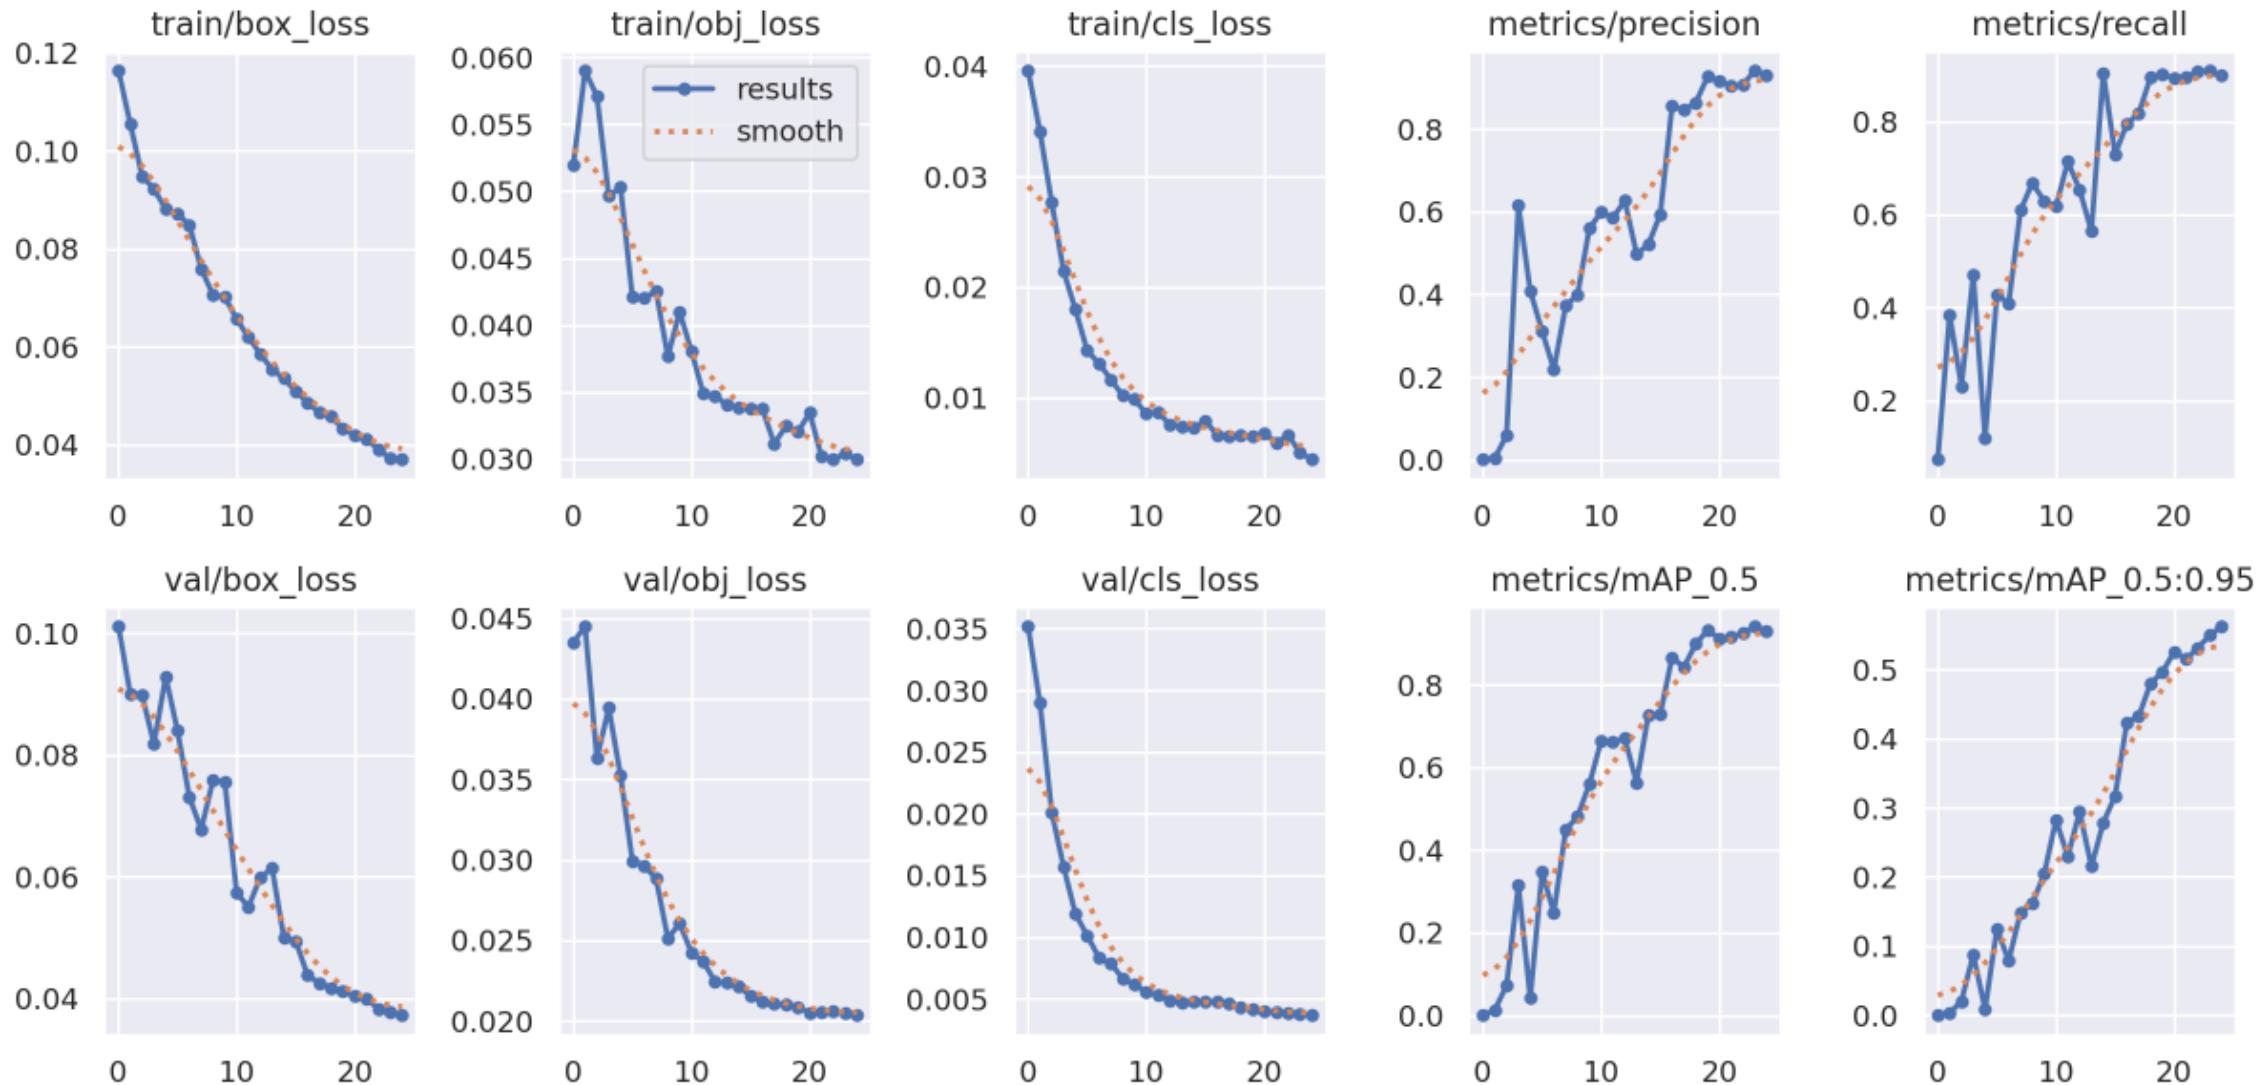
## Validation Loss:

- Validation loss also decreased, though at a slightly slower rate than the training loss, which indicates that the model is generalizing well to unseen data.
  - **Epoch 0: Validation Box Loss** = 0.10114
  - **Epoch 24: Validation Box Loss** = 0.037226

## Performance Metrics:

- The model shows a marked improvement in performance, as seen from the **precision**, **recall**, and **mAP** metrics:
  - **Epoch 0: Precision** = 0.073413, **Recall** = 0.0011638, **mAP@0.5** = 0.00083556, **mAP@0.5:0.95** = 0.00023351
  - **Epoch 24: Precision** = 0.92911, **Recall** = 0.0044069, **mAP@0.5** = 0.92868, **mAP@0.5:0.95** = 0.56218

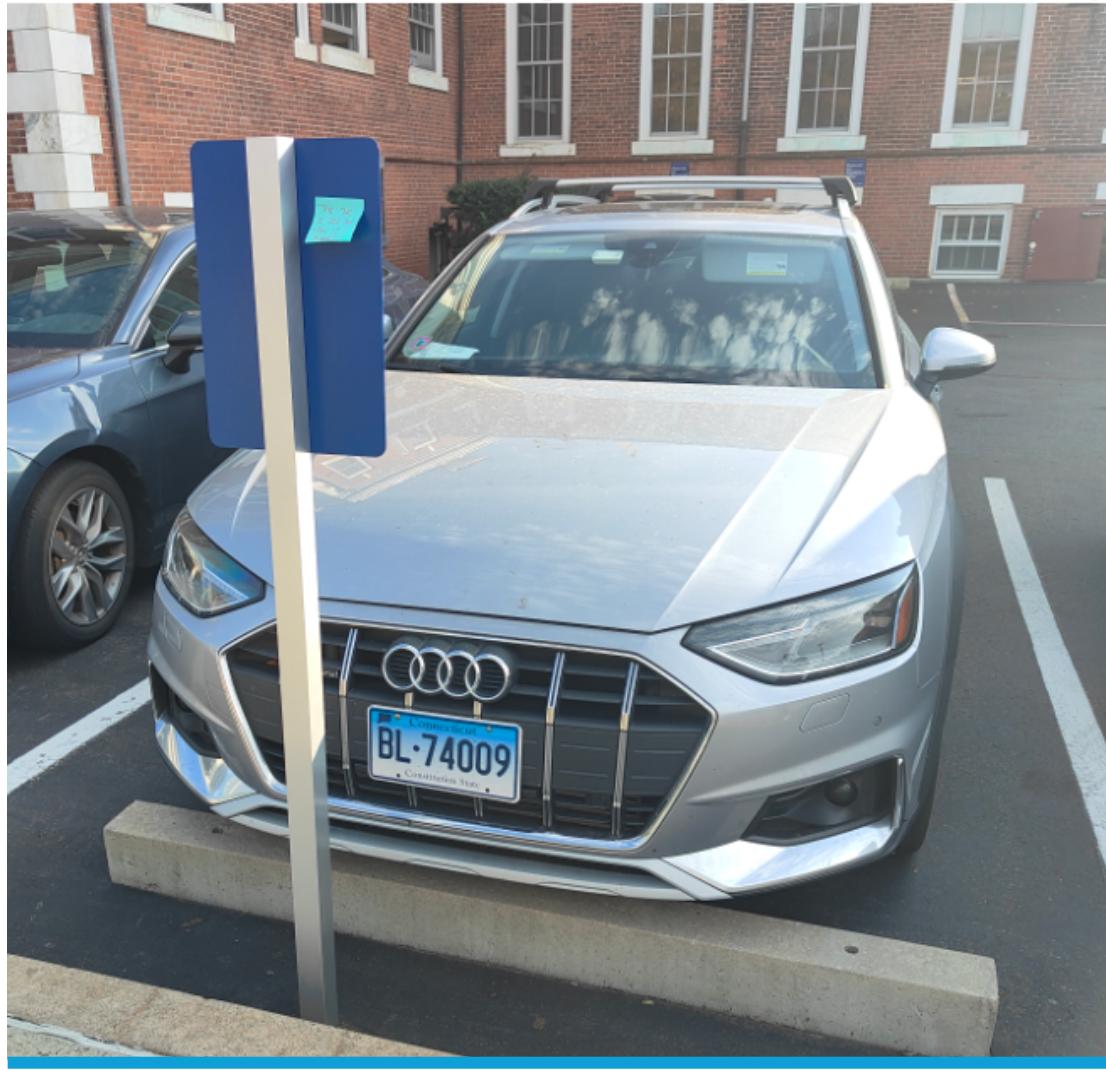
# Graphical Representation:



## BEFORE TRAINING:



## AFTER TRAINING:



# CONCLUSION:

**Before Training:** The model struggles with object detection and classification.

**After Training:** The model improves its object localization, classification, and confidence levels, indicating that it has learned the relevant patterns from the training data.